

MySQL Schema Best Practices

Mike Benshoof - Technical Account Manager, Percona



PERCONA
LIVE · USA
SANTA CLARA

Agenda

- .Introduction
- .Key things to consider and review
- .Tools/Queries to isolate issues
- .Common problems caused by schema design

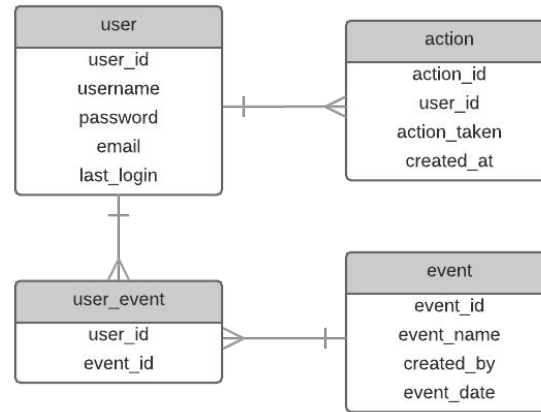
Introduction

•What comprises a “schema”?

- Formal definition of how data is organized
- Collection of tables, indexes, functions

Introduction - Sample Schema

Sample schema to highlight and refine:



Key Considerations

- Data Types
- Indexing

Data Types

•Key Principle

- Choose the smallest type you can

•Why does it matter?

- Disk space concerns
- Memory buffer allocation

Data Types

.Example:

- ~4 million rows, 4 columns, 3 indexes
- Identical data, one with bigint PK, one with unsigned int PK
- **Over 16% overhead!**

```
[root@plive-2017-demo plive_2017]# ls -alh action*.ibd
-rw-r-----. 1 mysql mysql 740M Apr 13 13:20 action_bigint.ibd
-rw-r-----. 1 mysql mysql 636M Apr 13 13:20 action.ibd
```

Data Types - Extra Considerations

•Numeric Types


- signed vs unsigned

•Date Types

- Timestamp for tracking (i.e. last update)
- Datetime for historical (i.e. birthday)

Indexing - High Level

Indexes serve to speed up data retrieval by using a secondary reference, just like the index in a normal book...



Indexing

•Key Principle

- Use fewest indexes needed to ensure optimal query performance

•Why does it matter?

- Indexes are needed for optimized query execution
- Each index has an additional cost
 - *Space (both disk and memory footprint)*
 - *Write performance hit (see InnoDB Change Buffer)*

Indexing - Extra Considerations

- PRIMARY KEY is appended to secondary indexes (InnoDB)
- Composite indexes move from left to right
 - Consider this to avoid duplicate indexes
 - Combine columns based on queries
 - Try to use more selective columns on the left
- Use EXPLAIN to verify which indexes are used

Profiling your Schema

•Quick Review: *pt-mysql-summary --databases=plive_2017*

```
# Schema #####
```

```
Database  Tables Views SPs Trigs Funcs   FKs Partn
plive_2017      5
```

```
Database  InnoDB
plive_2017      5
```

```
Database  BTREE
plive_2017      17
```

```
      i   v   t   b   c
      n   a   i   i   h
      t   r   m   g   a
           c   e   i   r
           h   s   n
           a   t   t
           r   a
               m
               p
```

```
Database  === === === === ===
plive_2017  6   6   4   1   2
```

Profiling your Schema

•Information_Schema

- Query directly
- Tables (estimates for InnoDB!)
- Indexes
- Stats_on_metadata (=0)

Finding Largest Schema

•Step 1:

- Find the largest schemas

```
SELECT table_schema, engine,  
count(*) tables,  
concat(round(sum(table_rows)/1000000,2),'M') rows,  
concat(round(sum(data_length)/(1024*1024*1024),2),'G') data,  
concat(round(sum(index_length)/(1024*1024*1024),2),'G') idx,  
concat(round(sum(data_length+index_length)/(1024*1024*1024),2),'G') total_size,  
round(sum(index_length)/sum(data_length),2) idxfrac  
FROM information_schema.TABLES  
WHERE table_schema not in ("information_schema", "mysql", "performance_schema", "sys")  
GROUP BY table_schema, engine  
ORDER BY sum(data_length+index_length) DESC LIMIT 20;
```

Finding Largest Schema

table_schema	engine	tables	rows	data	idx	total_size	idxfrac
plive_2017	InnoDB	5	10.81M	0.94G	0.84G	1.78G	0.89
test	InnoDB	2	0.79M	0.03G	0.00G	0.03G	0.00

Find Largest Tables in Schema

- Step 2:
 - Isolate the largest tables:

```
SELECT engine,  
table_name,  
concat(round(table_rows/1000000,2),'M') rows,  
concat(round(data_length/(1024*1024*1024),2),'G') data,  
concat(round(index_length/(1024*1024*1024),2),'G') idx,  
concat(round((data_length+index_length)/(1024*1024*1024),2),'G') total_size,  
round(index_length/data_length,2) idxfrac  
FROM information_schema.TABLES  
WHERE table_schema in ("plive_2017")  
ORDER BY data_length+index_length DESC LIMIT 10;
```


Find Largest Tables in Schema

engine	table_name	rows	data	idx	total_size	idxfrac
InnoDB	action	4.08M	0.30G	0.28G	0.59G	0.93
InnoDB	user_event	1.91M	0.22G	0.08G	0.30G	0.37
InnoDB	event	0.98M	0.11G	0.12G	0.23G	1.12
InnoDB	user	0.01M	0.00G	0.00G	0.00G	0.70

- engine, table_name
 - Just what you would think :)
- rows
 - Estimate of the number of rows
- data, idx, total_size
 - Estimate of the data/index size
- idxfrac
 - Ratio of index size : data size (ideally < 1)
 - Change ORDER BY to target these tables:
 - **ORDER BY idxfrac DESC LIMIT 10;**

What can you learn from this?

- (idxfrac) - which tables likely have:
 - Poor choices for PK data type
 - Excess indexes
- (rows/size) - which tables you should focus on for:
 - Archiving
 - Data retention
 - Partitioning

Finding duplicate keys...

The information_schema is a good way to potentially spot indexing issues at a high level...

Enter *pt-duplicate-key-checker* to quickly identify those problem indexes...

Finding duplicate keys...

```
[root@plive-2017-demo plive_2017]# pt-duplicate-key-checker --databases=plive_2017
# #####
# plive_2017.action
# #####

# idx_user is a left-prefix of idx_user_created
# Key definitions:
#   KEY `idx_user` (`user_id`),
#   KEY `idx_user_created` (`user_id`,`created_at`)
# Column types:
#   `user_id` int(11) default null
#   `created_at` timestamp not null default current_timestamp on update current_timestamp
# To remove this duplicate index, execute:
ALTER TABLE `plive_2017`.`action` DROP INDEX `idx_user`;

# #####
# Summary of indexes
# #####

# Size Duplicate Indexes    10205240
# Total Duplicate Indexes   1
# Total Indexes             13
```

Finding Unused Indexes

- Leverage the Performance Schema
 - Tracked since last restart

```
SELECT object_name, index_name
FROM
performance_schema.table_io_waits_summary_by_index_usage
WHERE index_name IS NOT NULL
AND count_star = 0
AND object_schema = "plive_2017"
ORDER BY object_name;
```

```
+-----+-----+
| object_name | index_name |
+-----+-----+
| action      | PRIMARY   |
| action      | idx_user  |
| action      | idx_created |
| action      | idx_user_created |
| event       | PRIMARY   |
| event       | idx_created_by |
| event       | idx_date   |
| user        | PRIMARY   |
| user        | idx_user   |
| user        | idx_email  |
| user        | idx_most_recent |
| user_event  | PRIMARY   |
| user_event  | idx_event  |
+-----+-----+
13 rows in set (0.00 sec)
```

Finding Unused Indexes

- Let's run some queries against our schema...

```
# Fetch all actions for user
SELECT *
FROM action
WHERE user_id = 104;
```

```
# Fetch all events a user is attending
SELECT username, event_id, event_name, event_date
FROM user JOIN user_event USING (user_id)
JOIN event USING (event_id)
WHERE user_id = 104;
```

```
# Find all users attending an event
SELECT username
FROM user
JOIN user_event USING (user_id)
WHERE event_id = "eb602e39-fe0d-11e5-a8b2-080027a5bc34";
```

Check Performance Schema...

```
SELECT object_name, index_name
FROM performance_schema.table_io_waits_summary_by_index_usage
WHERE index_name IS NOT NULL
AND count_star = 0
AND object_schema = "plive_2017"
ORDER BY object_name;
```

```
+-----+-----+
| object_name | index_name |
+-----+-----+
| action      | PRIMARY    |
| action      | idx_created |
| action      | idx_user_created |
| event       | idx_created_by |
| event       | idx_date    |
| user        | idx_user    |
| user        | idx_email   |
| user        | idx_most_recent |
+-----+-----+
8 rows in set (0.00 sec)
```

Great... now we can see the indexes that STILL haven't been touched!

Let's make some revisions...

- Drop duplicate index on **action** table
 - Found via pt-duplicate-key-checker
- Investigate data types / indexes in **event**
 - **Based on the $\text{idxfrac} > 1$**
 - I.e. indexes are larger than actual data

Drop duplicate keys

```
mysql> ALTER TABLE `plive_2017`.`action` DROP INDEX `idx_user`;
```

Before: -rw-r-----. 1 mysql mysql **636M** Apr 13 13:20 action.ibd

After: -rw-r-----. 1 mysql mysql **524M** Apr 13 13:53 action.ibd

That's an **18%** reduction in space!

Changing data types on event tables

```
mysql> show create table event\G
***** 1. row *****
      Table: event
Create Table: CREATE TABLE `event` (
  `event_id` char(36) NOT NULL DEFAULT '',
  `event_name` varchar(255) DEFAULT NULL,
  `created_by` int(10) unsigned DEFAULT NULL,
  `event_date` timestamp NOT NULL DEFAULT CURRENT_TIMESTAMP ON UPDATE CURRENT_TIMESTAMP,
  PRIMARY KEY (`event_id`),
  KEY `idx_created_by` (`created_by`),
  KEY `idx_date` (`event_date`)
) ENGINE=InnoDB DEFAULT CHARSET=latin1
```

```
mysql> show create table user_event\G
***** 1. row *****
      Table: user_event
Create Table: CREATE TABLE `user_event` (
  `user_id` int(10) unsigned NOT NULL DEFAULT '0',
  `event_id` char(36) NOT NULL DEFAULT '',
  PRIMARY KEY (`user_id`,`event_id`),
  KEY `idx_event` (`event_id`)
) ENGINE=InnoDB DEFAULT CHARSET=latin1
```

Changing data types on event tables

Transitional tables:

```
CREATE TABLE `event_new` (  
  `id` int(10) unsigned NOT NULL AUTO_INCREMENT,  
  `event_id` char(36) NOT NULL DEFAULT '',  
  `event_name` varchar(255) DEFAULT NULL,  
  `created_by` int(10) unsigned DEFAULT NULL,  
  `event_date` timestamp NOT NULL DEFAULT CURRENT_TIMESTAMP ON UPDATE CURRENT_TIMESTAMP,  
  PRIMARY KEY (`id`),  
  KEY `idx_created_by` (`created_by`),  
  KEY `idx_date` (`event_date`)  
) ENGINE=InnoDB
```

```
CREATE TABLE `user_event_new` (  
  `user_id` int(10) unsigned NOT NULL DEFAULT '0',  
  `event_id` int(10) unsigned NOT NULL DEFAULT '0',  
  PRIMARY KEY (`user_id`, `event_id`),  
  KEY `event_id` (`event_id`)  
) ENGINE=InnoDB
```



PERCONA

Changing data types on event tables

Populating Transitional Tables & Cleanup:

```
# Generate an integer PK for each row
```

```
mysql> INSERT INTO event_new SELECT NULL, event.* FROM event;
```

```
Query OK, 1048576 rows affected (12.25 sec)
```

```
Records: 1048576 Duplicates: 0 Warnings: 0
```

```
# Replace each char-based event_id with the new integer
```

```
mysql> INSERT INTO user_event_new SELECT user_id, event_new.id FROM user_event JOIN event_new USING  
(event_id);
```

```
Query OK, 2048431 rows affected (22.20 sec)
```

```
Records: 2048431 Duplicates: 0 Warnings: 0
```

```
# Remove the old character event_id field
```

```
mysql> ALTER TABLE event_new DROP COLUMN event_id;
```

```
Query OK, 0 rows affected (19.22 sec)
```

```
Records: 0 Duplicates: 0 Warnings: 0
```

Changing data types on event tables

Review the Impact:

```
-rw-r-----. 1 mysql mysql 252M Apr 13 13:17 event.ibd  
-rw-r-----. 1 mysql mysql 132M Apr 13 14:04 event_new.ibd  
  
-rw-r-----. 1 mysql mysql 332M Apr 13 13:18 user_event.ibd  
-rw-r-----. 1 mysql mysql 172M Apr 13 14:02 user_event_new.ibd
```

That is a 48% reduction in space!



PERCONA

Schema breakdown - revisited...

engine	table_name	rows	data	idx	total_size	idxfrac
InnoDB	action	4.17M	0.35G	0.14G	0.50G	0.41
InnoDB	user_event	1.91M	0.22G	0.08G	0.30G	0.37
InnoDB	event	0.98M	0.11G	0.12G	0.23G	1.12
InnoDB	user_event_new	1.91M	0.12G	0.03G	0.15G	0.21
InnoDB	event_new	1.04M	0.09G	0.03G	0.12G	0.38
InnoDB	user	0.01M	0.00G	0.00G	0.00G	0.70

Common Issues

- Wasted resources
 - Memory (buffer pool, per-session buffers, etc)
 - Disk
- Sub-optimal performance
 - Primarily write operations
- Additional overhead for optimizer
 - More execution paths to analyze
 - Higher likelihood for wrong path to be chosen

The times, they are a changin'...

- Schema review is an iterative process, not a one time event!
- Index utilization isn't remotely static
 - Periodically refresh P_S, especially after code and/or index changes
 - `TRUNCATE TABLE performance_schema.table_io_waits_summary_by_index_usage`
- Track size over time for trending
 - Help to answer the common question: “When do we need more disk space?”
 - Determine the frequency of archiving/pruning

MySQL Schema Review 101

Questions?

MySQL Schema Review 101

Thanks!